

Introduction to Deep Learning

Pablo Lamata
Eric Kerfoot
Esther Puyol

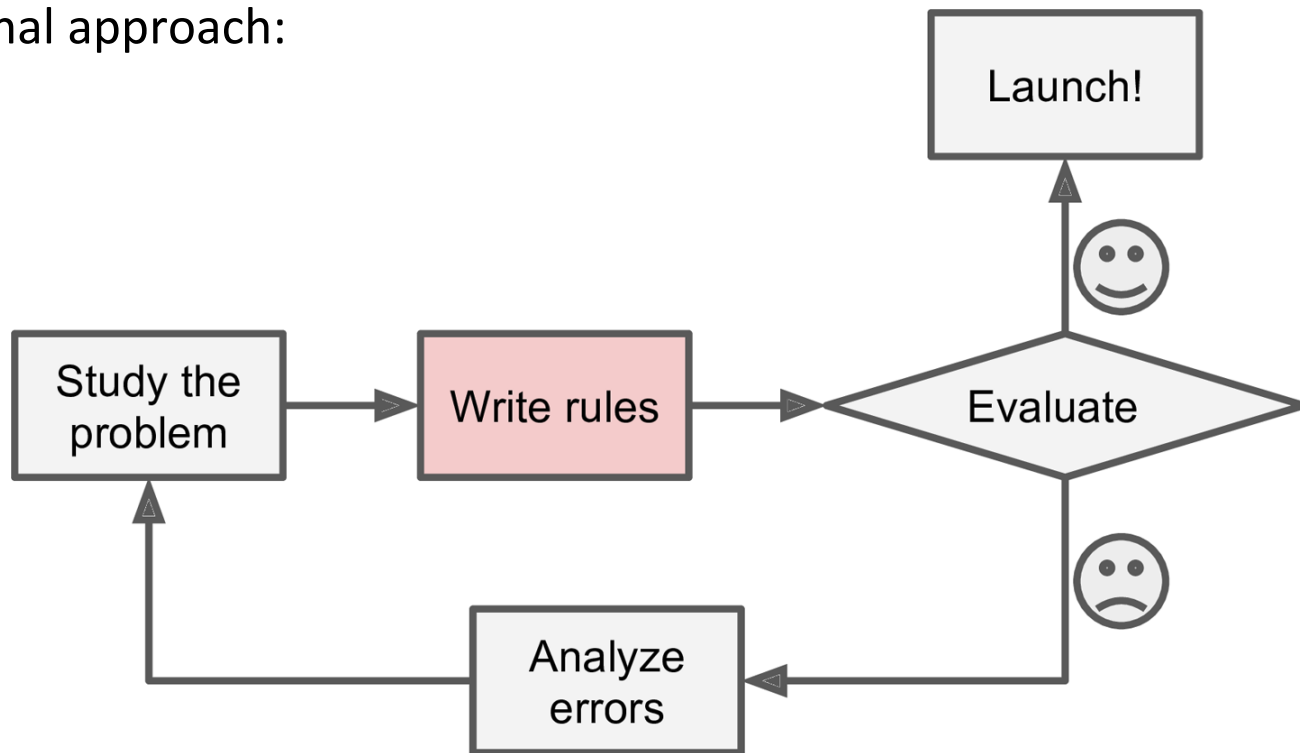
School of Biomedical Engineering & Imaging Sciences
King's College London



What is AI?

The science (and art) of programming computers so they can *“learn from and make predictions on data”*

Traditional approach:

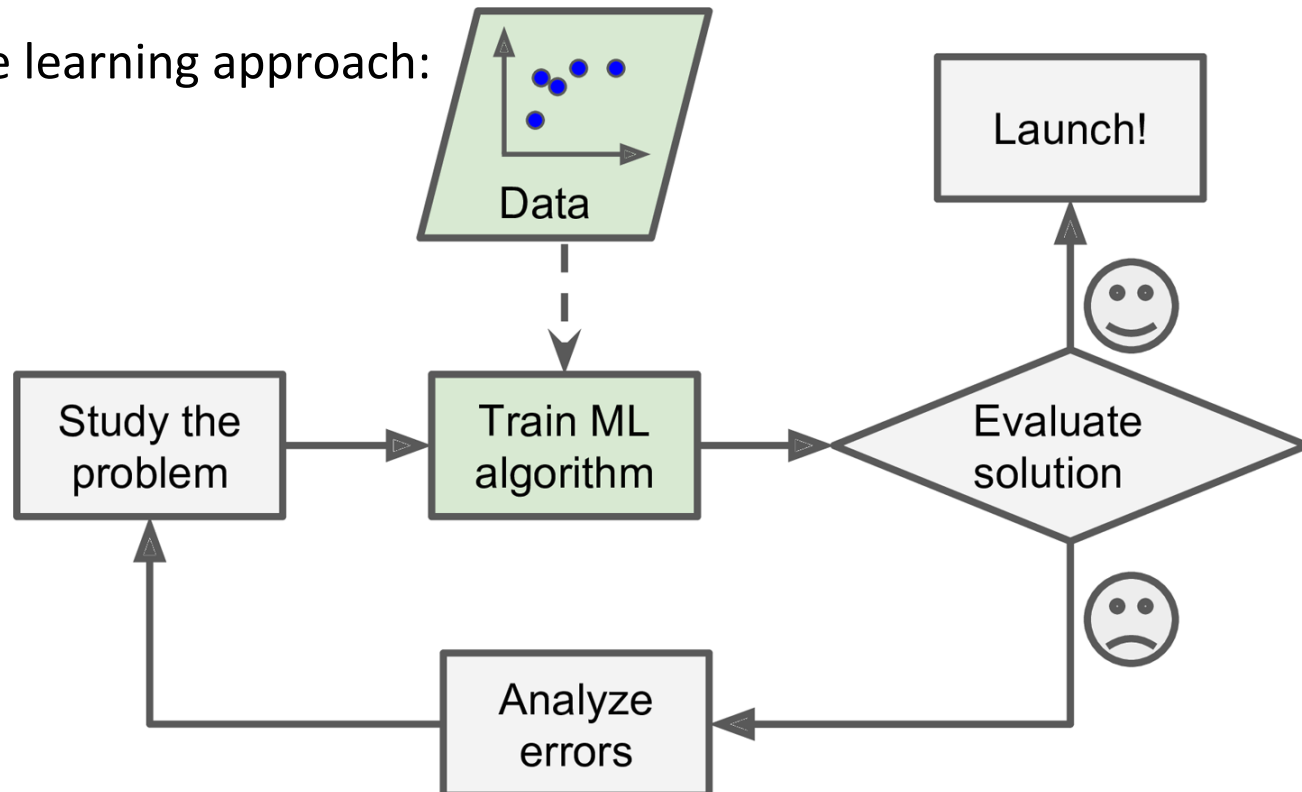




What is machine learning?

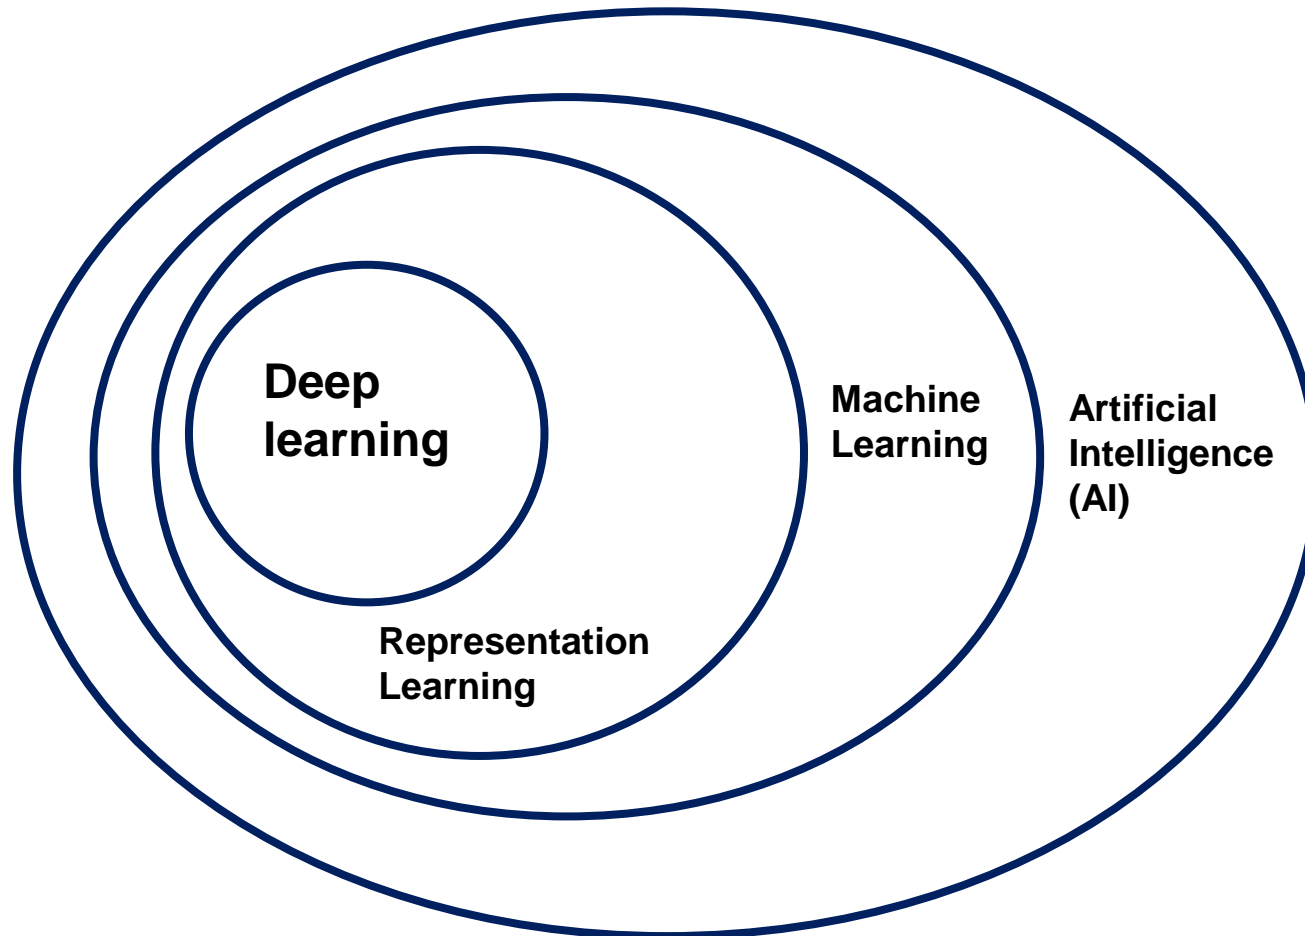
The science (and art) of programming computers so they can *“learn from and make predictions on data”*

Machine learning approach:





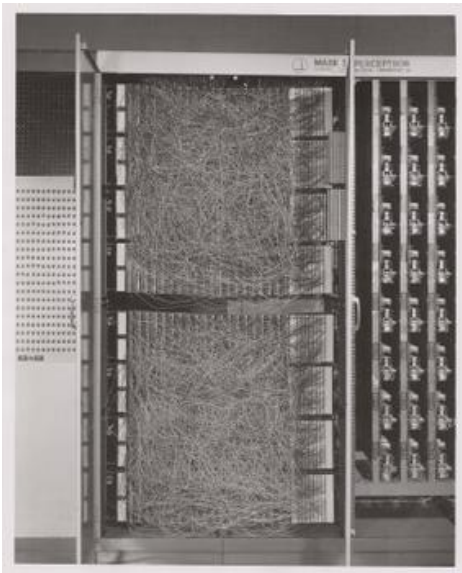
AI, Machine Learning and Deep Learning:





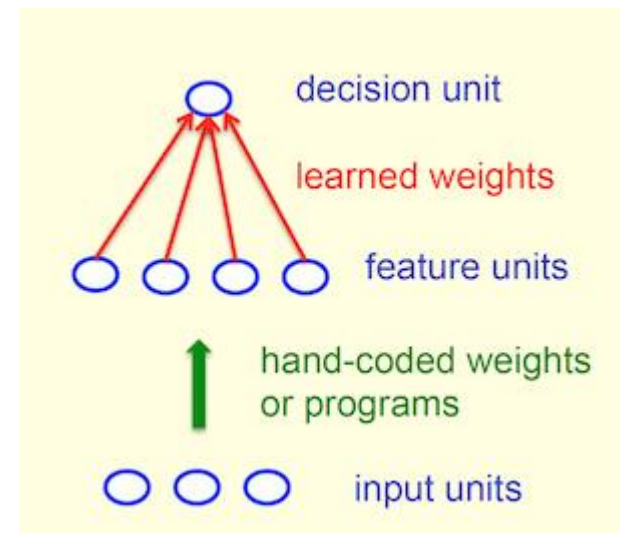
Perceptron

The first neural network (Frank Rosenblatt, 1957)



$$f(x) = \begin{cases} 1 & \text{if } \sum_i w_i \cdot x_i > b \\ 0 & \text{otherwise} \end{cases}$$

"Mark 1 perceptron" - machine designed for image recognition: it had an array of 400 [photocells](#), randomly connected to the "neurons". Weights were encoded in [potentiometers](#), and weight updates during learning were performed by electric motors





Application of AI

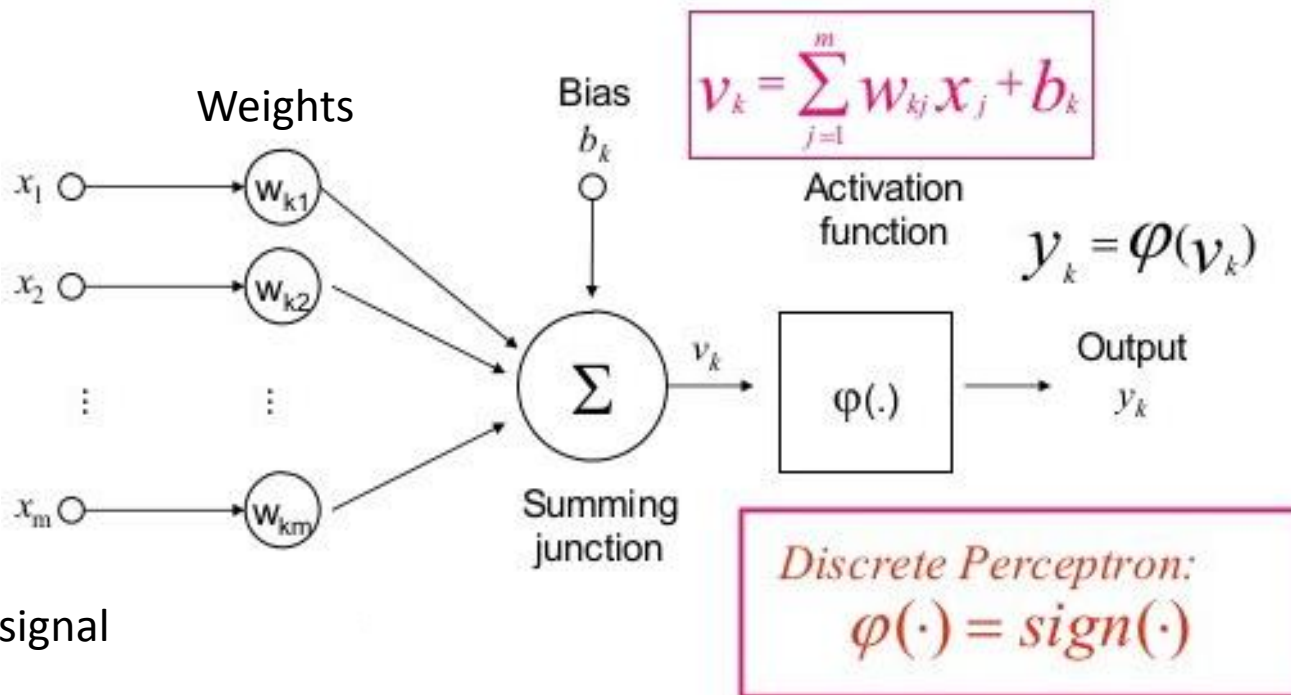
- Optical character recognition
- Spam filtering
- Speech recognition
- Face recognition
- Object recognition
- Robotics
- Recommending products based on what customers bought before
- Detect fraudulent transactions





A single-layer perceptron

A single-layer perceptron looks as follows:



*This and the following slides follow the example on <https://hackernoon.com/a-hands-on-introduction-to-neural-networks-6a03afb468b1>



A single-layer perceptron

Input:

- Each input to the neuron (x_1, x_2, \dots, x_n) is known as a **feature**

Weights:

- Each feature is weighted with a number to represent the strength of that input ($w_{k1}, w_{k2}, \dots, w_{km}$).

Activation function:

- Calculate weighted sum of inputs (v_k), pass through an **activation function** and threshold result y_k to 0 or 1

A single-layer perceptron can be trained as follows:

1. Ask the neuron* to classify a sample (**forward pass**)
2. Update the neuron's weights based on how wrong the prediction is.
3. Repeat for a set number of times (=epochs).



Activation function

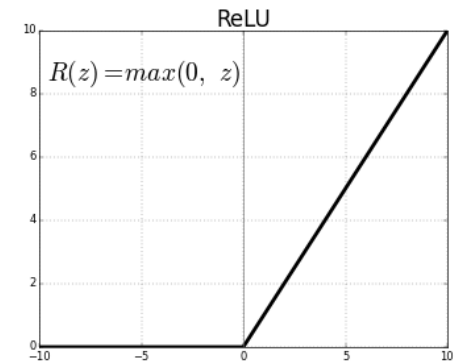
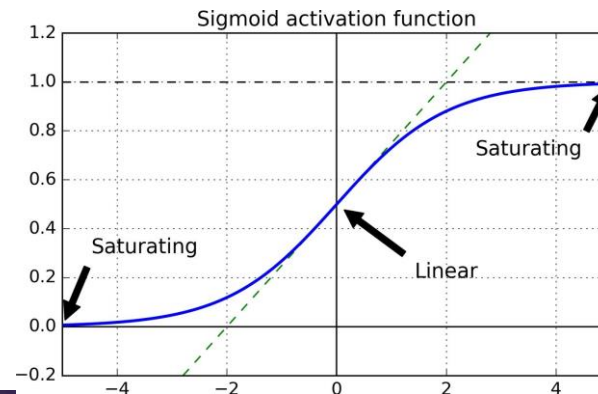
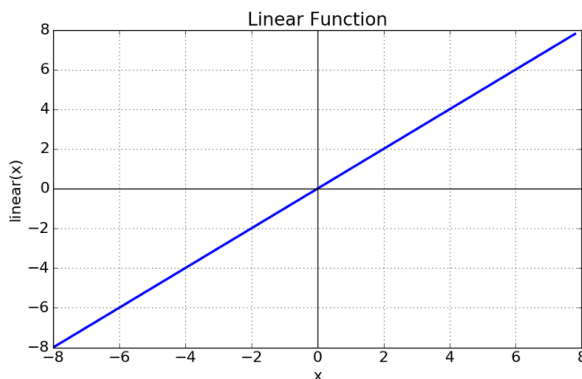
Added to the output end of any neural network

- Can be regarded as a **Transfer Function**
- Can also be attached in between two Neural Networks

Used to determine whether the output of a neural network is 'yes' or 'no' (or something in between).

- Maps the resulting values in between 0 to 1 or -1 to 1 (depending upon the activation function)

We distinguish between (piecewise) linear and nonlinear activation functions





Training a perceptron

What else do we need?

- We need **an error function (cost function or loss function)**:
 - The cost function must be able to be written as an average over cost functions E_i for individual training examples x_i :

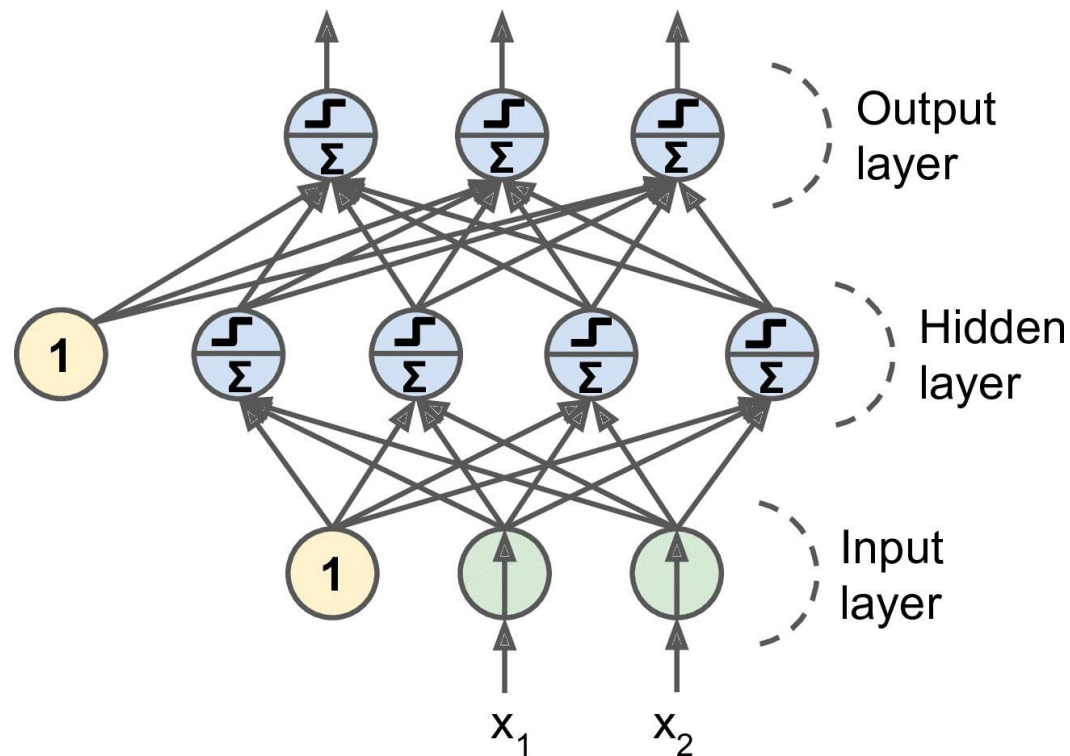
$$E = \frac{1}{N} \sum_{i=1}^N E_i$$

- This allows us to compute the gradient (with respect to weights and biases) for a single training example, and run e.g. Gradient Descent
- Examples of loss function: L1 norm $\rightarrow E = \sum_i |E_i|$
- We need an **optimisation method** (e.g. gradient descent)
- Our activation function should be **differentiable**



From single-layer to multi-layer perceptron

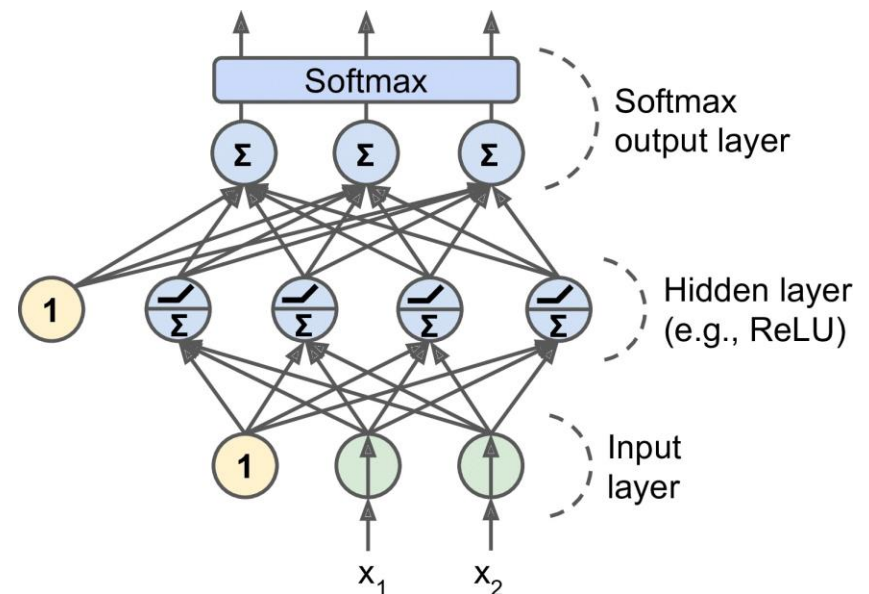
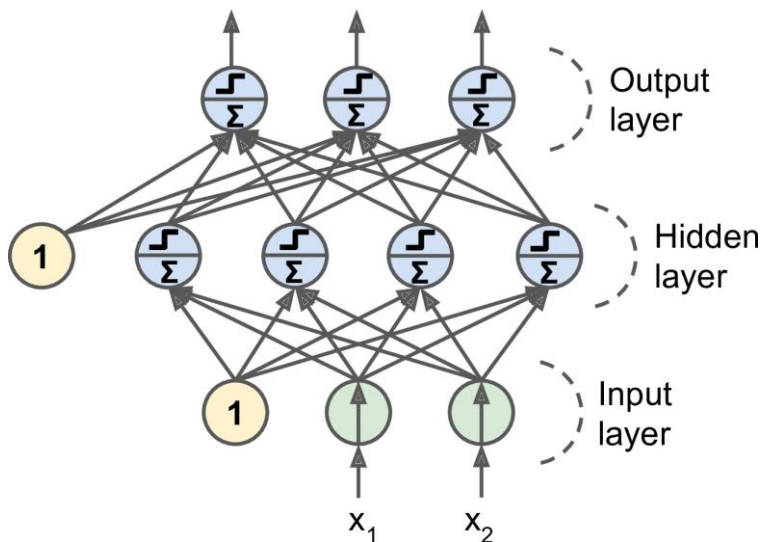
A more complex MLP is shown below (still only 1 hidden layer):





Multi-layer perceptron

We have stacked multiple perceptrons to generate hidden layers:



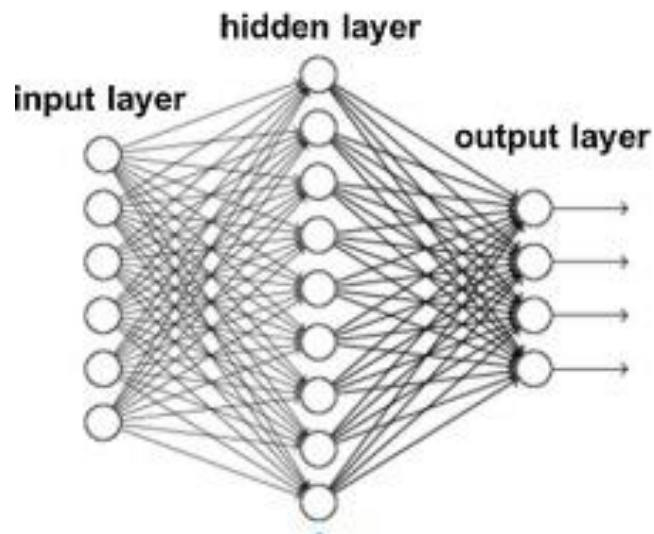
If we have more than one hidden layer, the neural network is considered to be “deep” and we move into **deep learning**



From shallow to deep fully connected networks

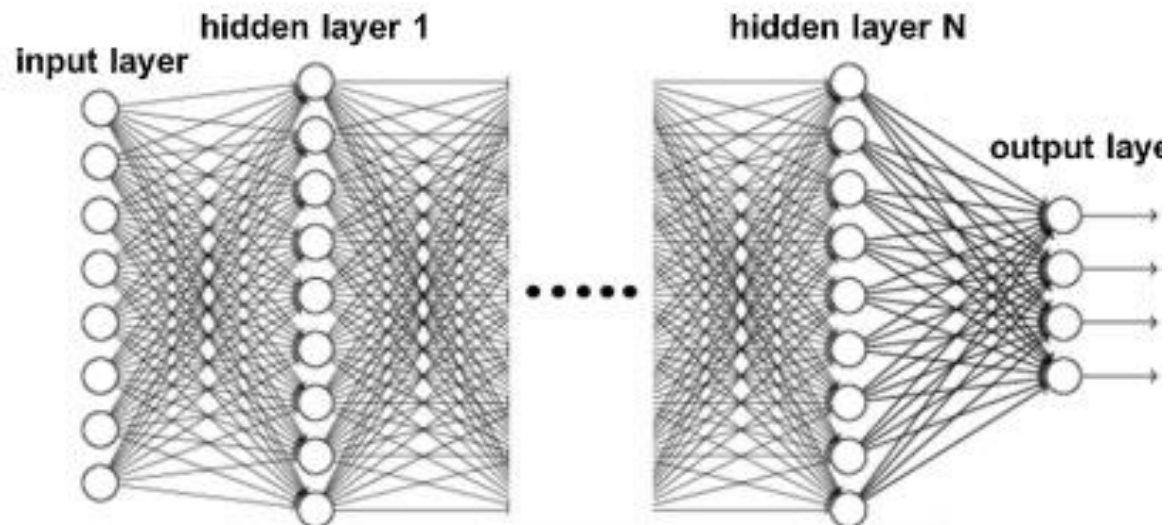
Compare this to a deep fully-connected network with N hidden layers:

Shallow neural network



$6 \times 9 = 54$ connections to first layer
 $9 \times 4 = 36$ connections to output layer
 $54 + 36 = 90$ connections in total

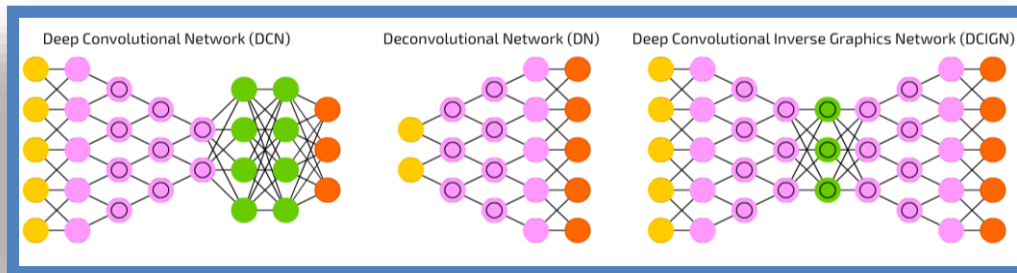
Deep neural network



$8 \times 9 = 72$ from input layer, $9 \times 4 = 36$ to output layer
 9×9 connections between each hidden layer, ie $(9 \times 9)^{(N-2)}$
 $8 \times 9 + (N-2) \times 81 + 9 \times 4$
 E.g. **108** (N=2), **189** (N=3), **270** (N=3)



Convolutional neural network (CNN)



Convolutional neural networks

Input Cell

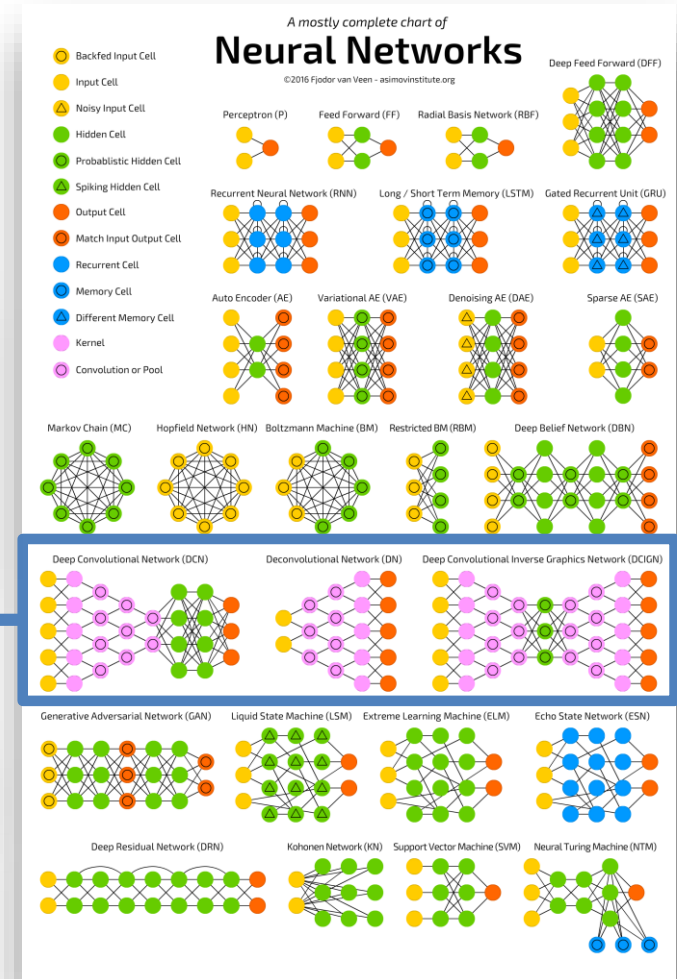
Hidden Cell

Probabilistic Hidden Cell

Output Cell

Kernel

Convolution or Pool



[Src: towardsdatascience.com](http://towardsdatascience.com)



Convolutional Neural Networks (CNNs)

CNNs have several important **building blocks**:

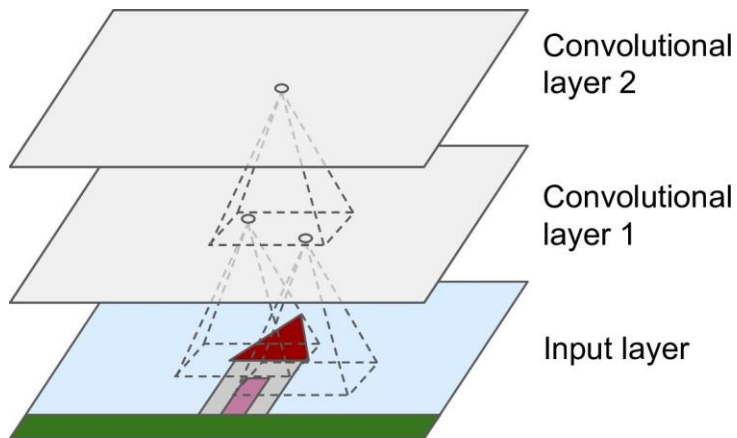
1. **2D (or 3D) input layer**
2. **Convolutional layer**
 - Neurons in the first convolutional layer are not connected to every single pixel, **but only to pixels in their receptive fields**
 - Neurons in the second convolutional layer are only connected to neurons within a small rectangular region in the first layer
3. **Pooling layer**
 - Goal is to **subsample** the input image to reduce computational load, memory usage, numbers of parameters (limits risk of overfitting)
4. **Fully-connected output layer**
 - This is a **regular feed-forward network** which produces final output prediction
 - E.g. **softmax layer** that outputs estimated class probabilities



Convolutional layer

Convolutional layers allow the network to:

- Concentrate on low-level features in the first hidden layer
- Assemble them to higher-level features in the next hidden layer



1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

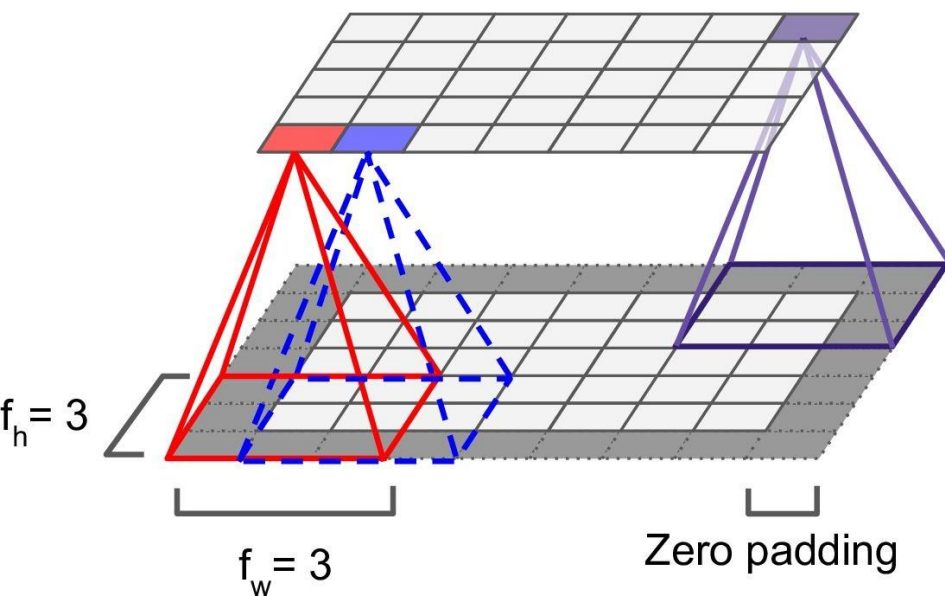
4		

Convolved
Feature

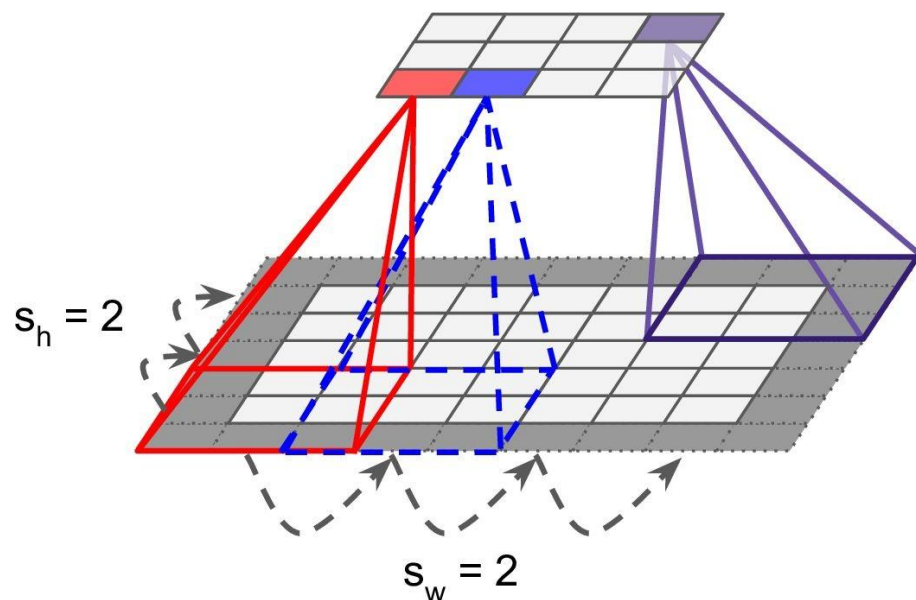


Convolutional layer

Need to apply **zero-padding** at each layer and also apply a stride for further dimensionality reduction:



Connections between layers and zero padding

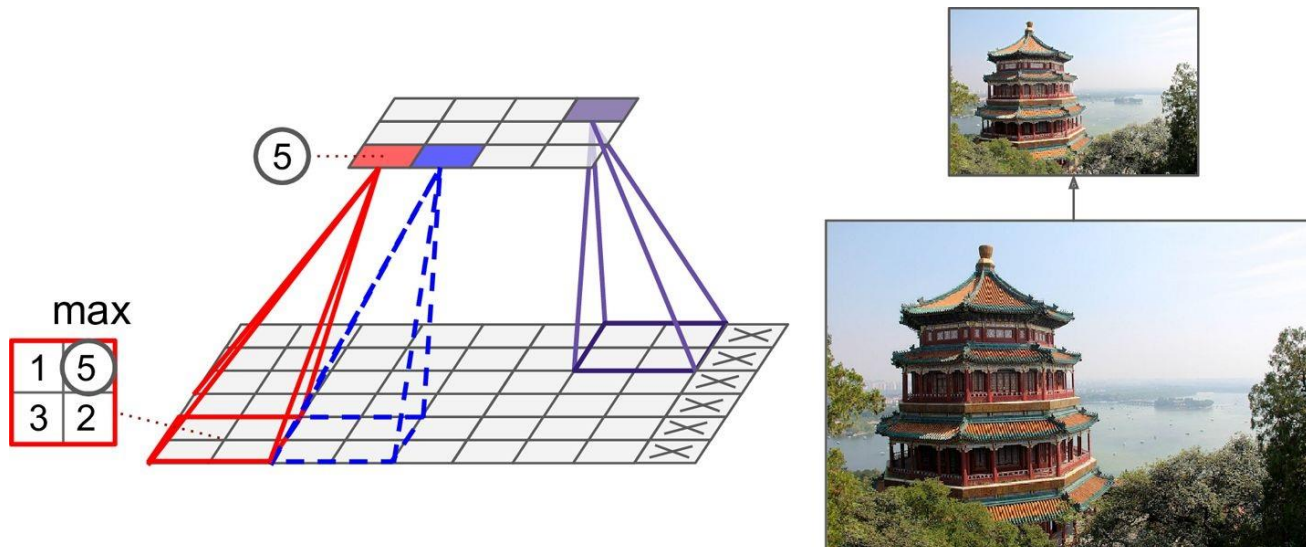


Reducing dimensionality using a stride



Pooling layer

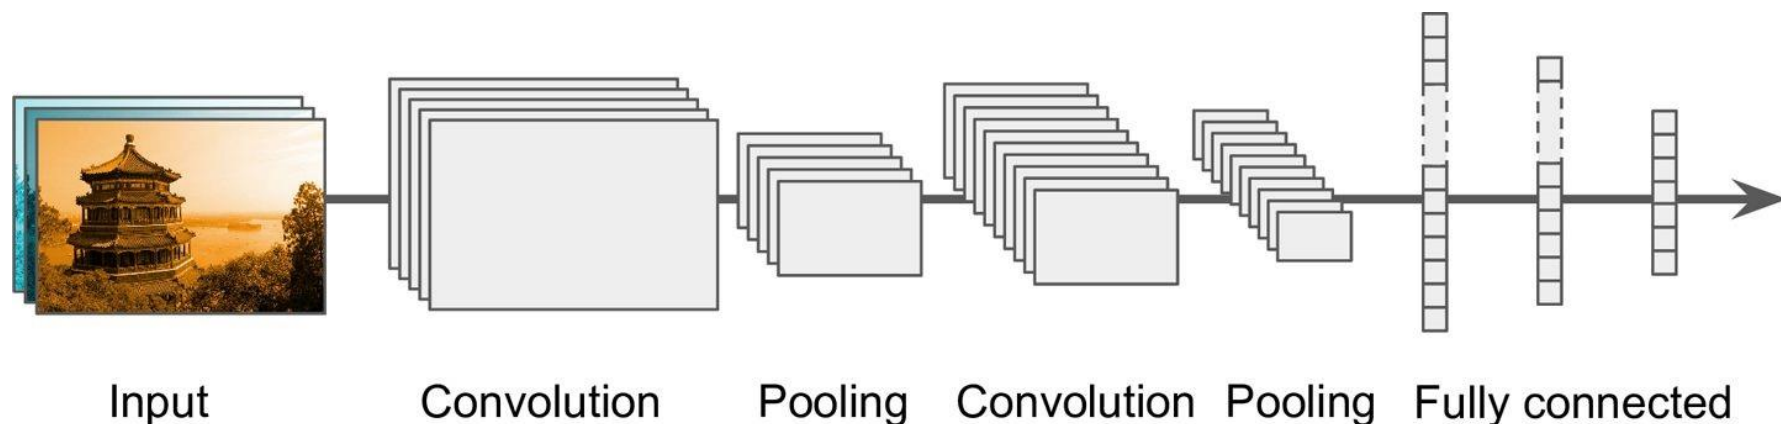
- Goal is to **subsample** the input image to reduce computational load, memory usage, numbers of parameters (limits risk of overfitting)
- Each neuron in pooling layer is connected to the outputs of limited number of neurons in previous layer, again located within a small rectangular receptive field
- Most common type is **max pooling**





Convolutional Neural Networks (CNNs)

Putting it all together:



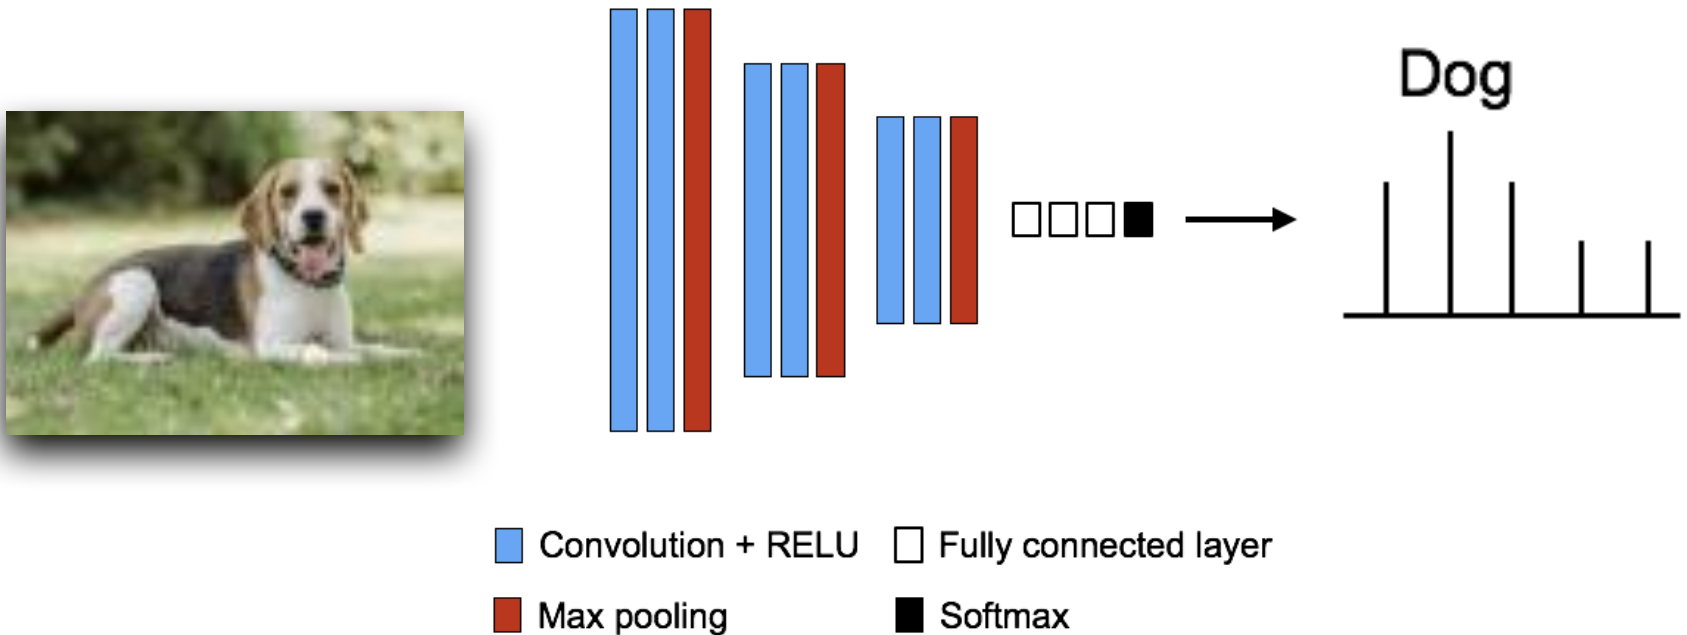
In this case, 3-channel
RGB image)

Many variants of CNN exist and have been boosted by the
advent of **ImageNet** in 2010



CNNs for image classification

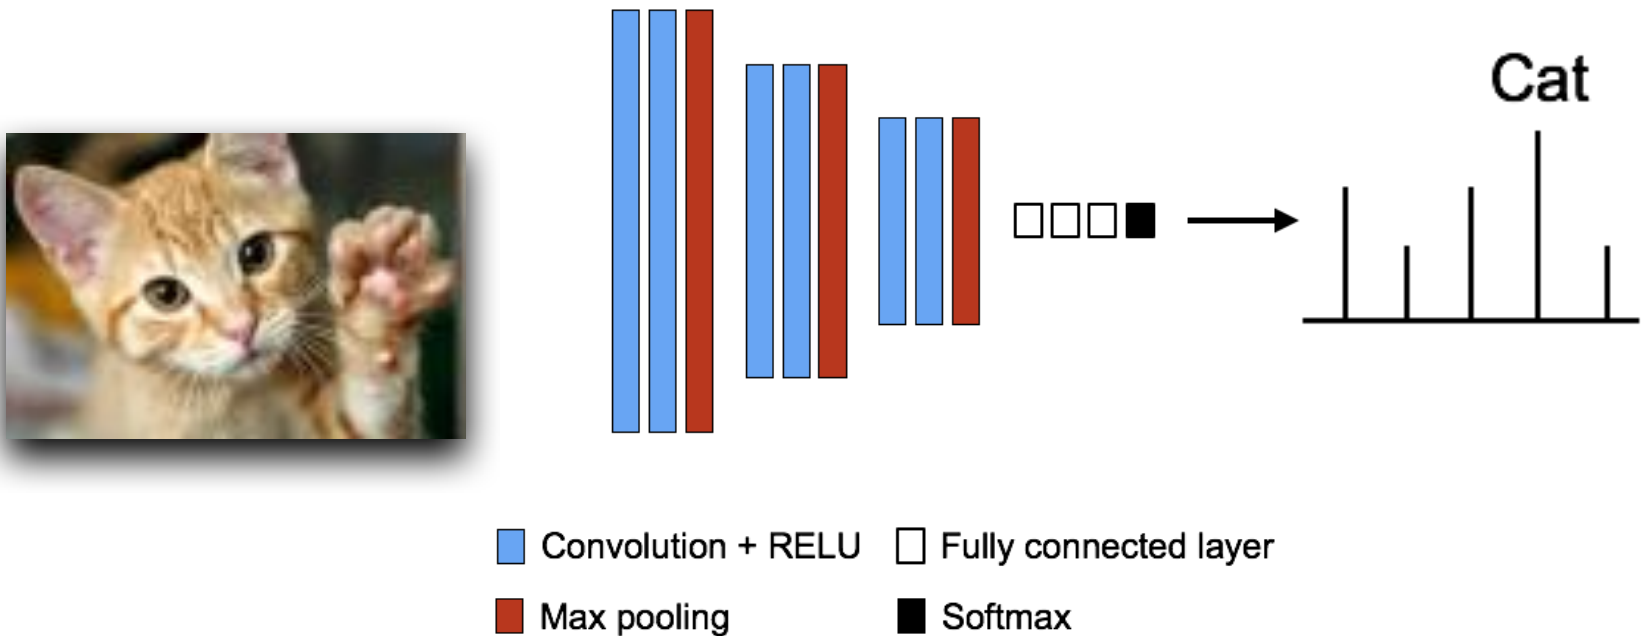
First rewind and look at a simple CNN applied to real image classification





CNNs for image classification

First rewind and look at a simple CNN applied to real image classification

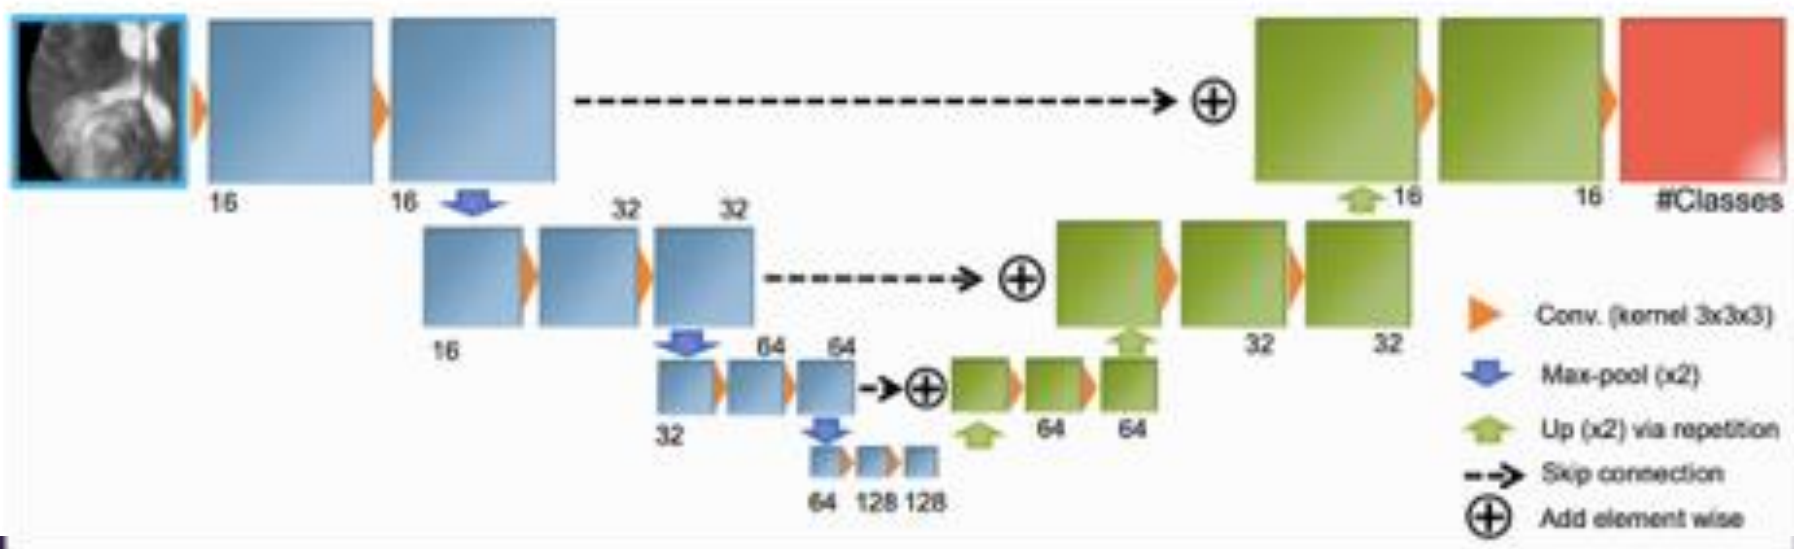




CNNs for image segmentation

One major focus in computer vision and medical imaging is on **image segmentation**

- Challenging in medical imaging due to variability in patient anatomy & pathology, patient pose and motion, image artefacts
- U-net is the most common used network , which was proposed by Ronneberger (Google DeepMind) in 2015:





Bits & Bobs

All of the networks so far (deep neural networks, with or without convolutional layers), need to be carefully designed and trained:

Choice of:

- Number of hidden layers and neurons, stacking (AE)
- Loss function, activation function, learning rate, epochs
- **Optimisation methods**
- **Regularisation methods**
- **Transfer learning**



Optimisation

Choice of faster gradient-based optimisation methods for use with backpropagation:

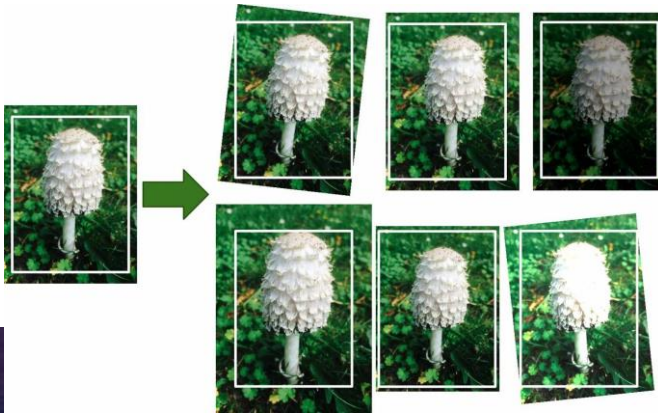
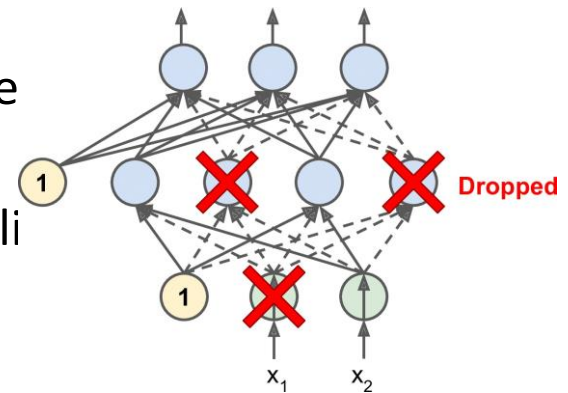
1. **Gradient Descent**
2. **Momentum optimisation**
3. **Nesterov Accelerated Gradient**
4. **AdaGrad**
5. **RMSPProp**
6. **Adam Optimiser**



Regularisation

To avoid **overfitting**, we can do the following:

1. **Early stopping:** Interrupt training when its performance on the validation set starts dropping
2. **l_1 and l_2 regularization:** Add a regularization term in the cost function.
3. **Dropout:** At every training step, every neuron (input or hidden) has a probability p of being temporarily “dropped out”
4. **Data augmentation:** Generate new training instance artificially boosting the size of the training set.
E.g. you can rotate, shift (translate), resize (scale), fli





Summary

We can also exploit existing network architectures or parts thereof by:

- Stacking** different types of layers, such as convolutional and pooling layers – it's very modular

- Creating **ensemble** architectures by aggregating results from different, separately trained architectures

- Using pre-trained networks (e.g. U-net) for similar applications (**transfer learning**)

Questions?

Slides adapted from the Machine Learning for
Biomedical Application course from King's College of
London